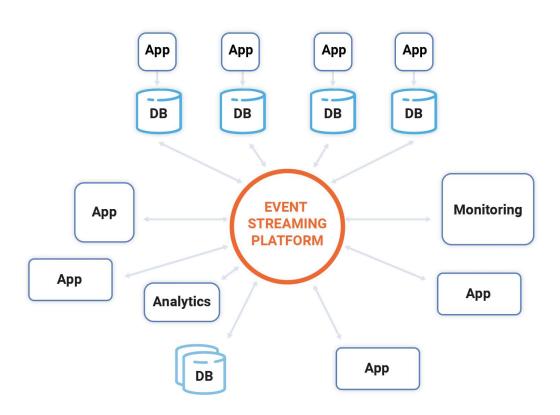
Distributed and Parallel Computing

Trong-Hop Do

Kafka – A distributed event streaming flatform



What is event streaming?

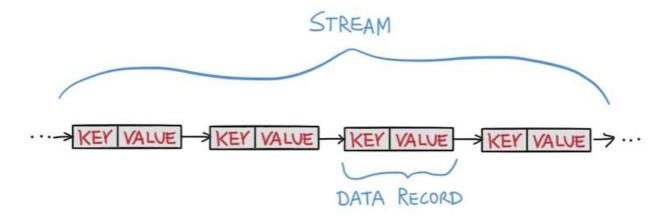


What can I use event streaming for?

- To process payments and financial transactions in real-time, such as in stock exchanges, banks, and insurances.
- To track and monitor cars, trucks, fleets, and shipments in real-time, such as in logistics and the automotive industry.
- To continuously capture and analyze sensor data from IoT devices or other equipment, such as in factories and wind parks.
- To collect and immediately react to customer interactions and orders, such as in retail, the hotel and travel industry, and mobile applications.
- To monitor patients in hospital care and predict changes in condition to ensure timely treatment in emergencies.
- To connect, store, and make available data produced by different divisions of a company.
- To serve as the foundation for data platforms, event-driven architectures, and microservices.

What is a stream?

- Think of a stream as an unbounded, continuous real-time flow of records
 - You don't need to explicitly request new records, you just receive them
- Records are key-value pairs



The Shift to Event-driven Systems has Already Begun...

From a static snapshot...

...to a continuous stream of events



Occasional call to a friend



A constant feed about the activities of all your friends



Daily news reports









Real time news feeds, accessible online anytime, anywhere

This leads us to...



Single platform to connect everyone to every event



Real-time stream of events



All events stored for historical view

Successful Digital Businesses are Inherently Event-driven

Born cloud-native...

Traditional ones that adapt...



Social Networks
Enabling Event
Sharing



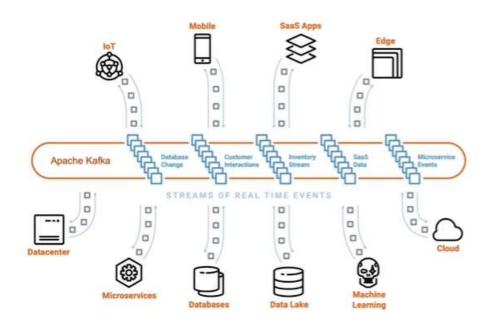
Newspaper Provide a single Source of Truth



Streaming Provider
On-demand Digital
Content



Credit Card Payments
Microservices
Architecture



Apache Kafka®: the Defacto Standard for Real-Time Event Streaming

- Global-scale
- Real-time
- Persistent Storage
- Stream Processing

Thousands of Companies Worldwide trust Kafka for their Journey towards "**Event-driven**"

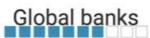






6 of top 10





7 of top 10







9 of top 10

Real-time Fraud Detection

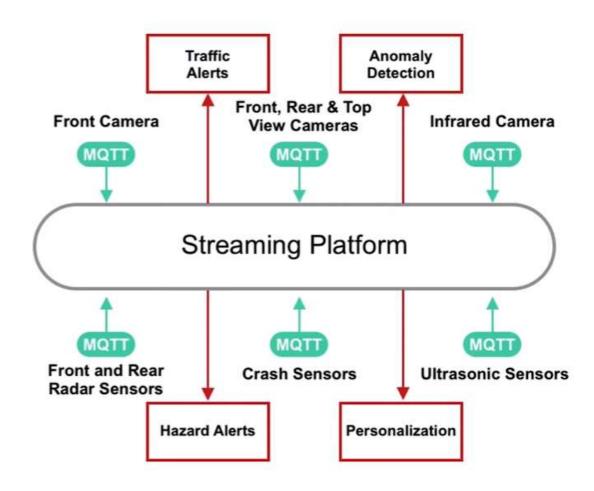


- Act in real-time
- Detect fraud
- Minimize risk
- Improve customer experience

Automotive



The Future of the Automotive Industry is a Real Time Data Cluster



Real-time e-Commerce



Rewards Program

- Onboarding new merchants faster
- Increased speed at which mobile applications are delivered to customers
- Enabled a full 360 view of customers
- Enhanced performance and monitoring
- Projected savings of millions of dollars

Health Care

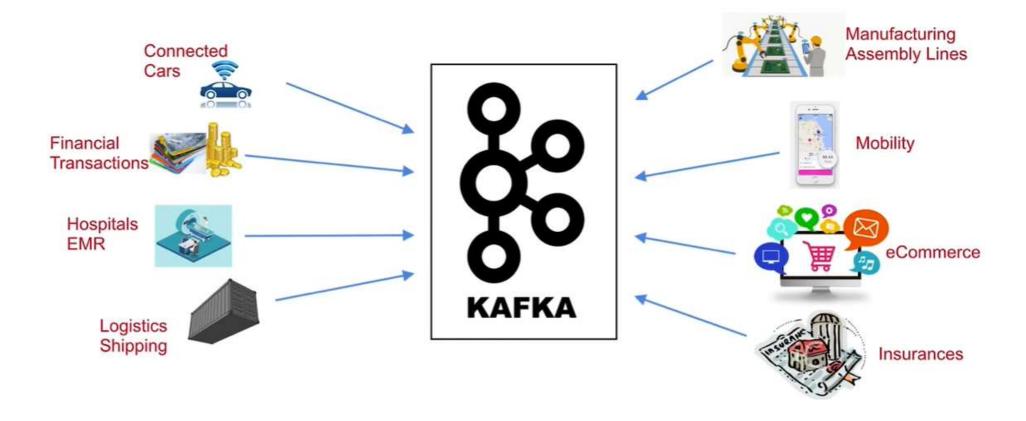


Microservices

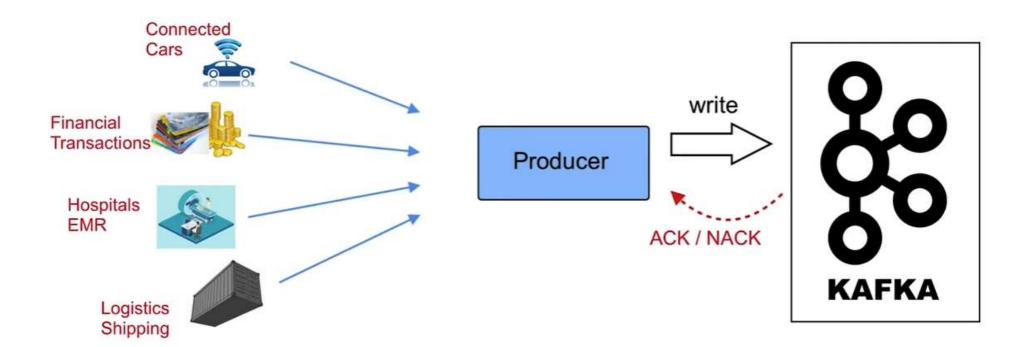


Internet of Things

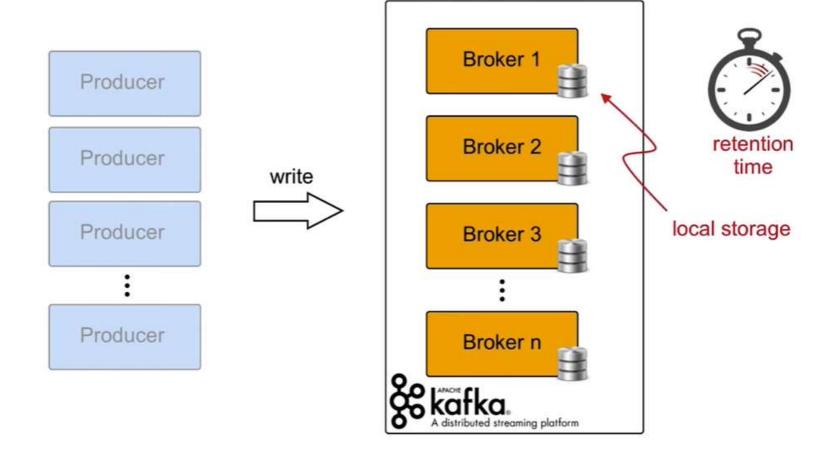
The World Produces Data



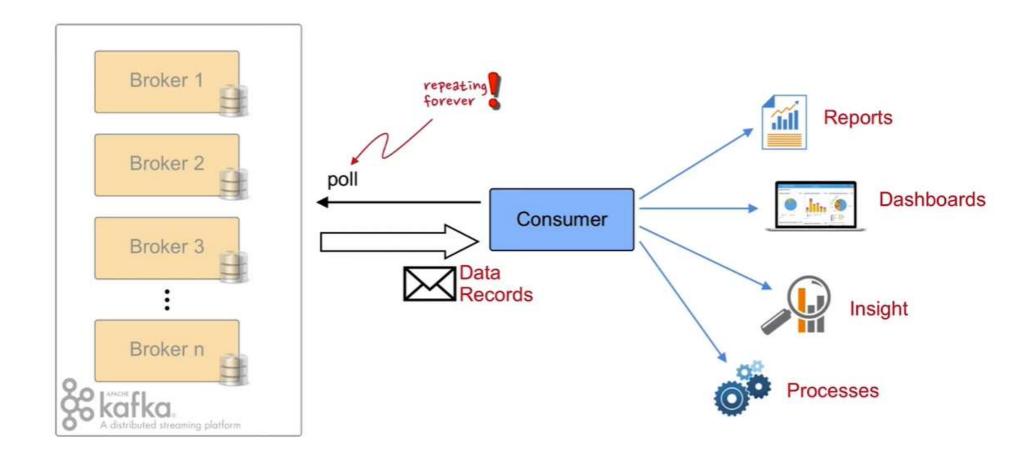
Producers



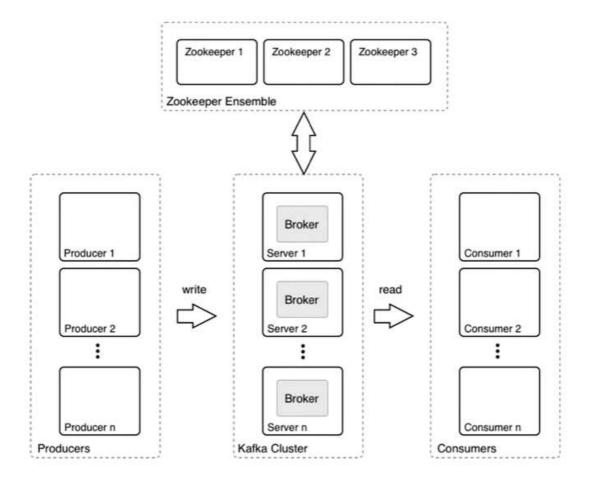
Kafka Brokers



Consumers

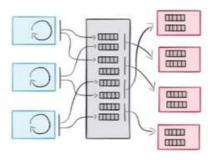


Architecture

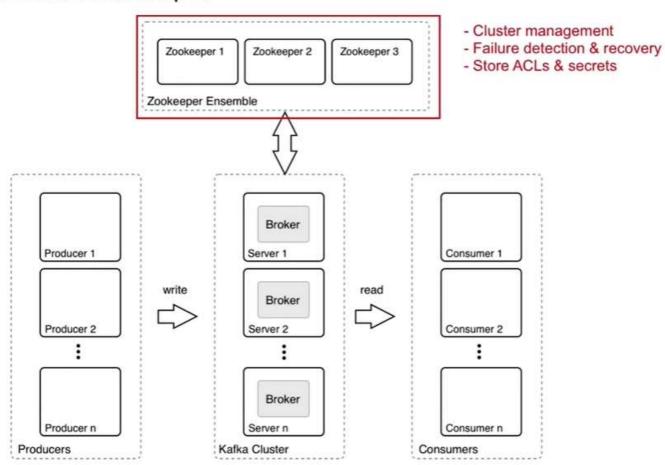


Decoupling Producers and Consumers

- Producers and Consumers are decoupled
- Slow Consumers do not affect Producers
- Add Consumers without affecting Producers
- Failure of Consumer does not affect System

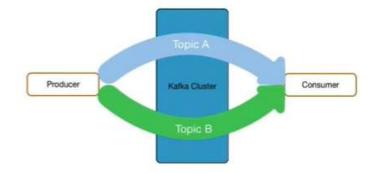


How Kafka Uses ZooKeeper

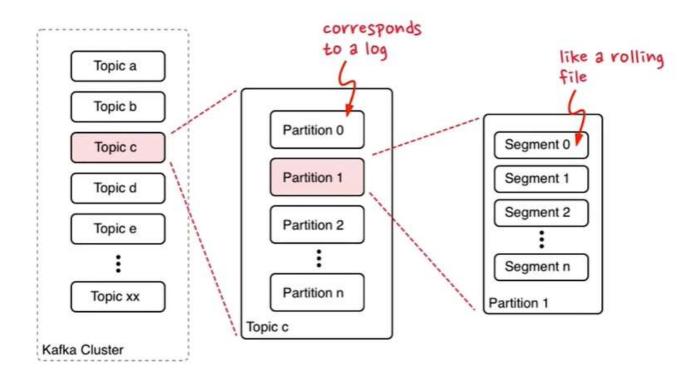


Topics

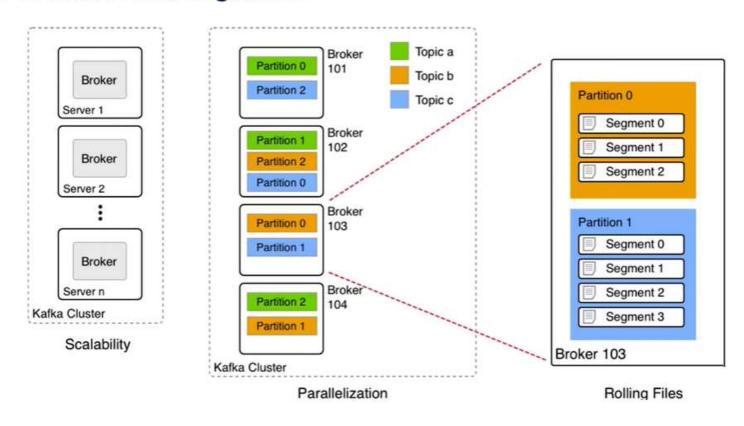
- Topics: Streams of "related" Messages in Kafka
 - Is a Logical Representation
 - Categorizes Messages into Groups
- Developers define Topics
- Producer ←→ Topic: N to N Relation
- Unlimited Number of Topics



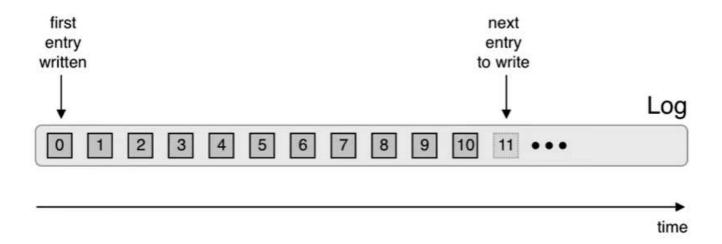
Topics, Partitions and Segments



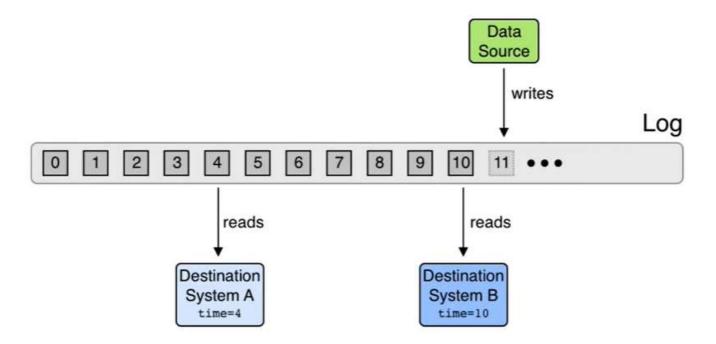
Topics, Partitions and Segments



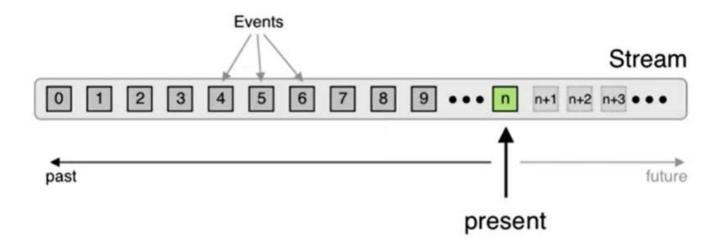
The Log



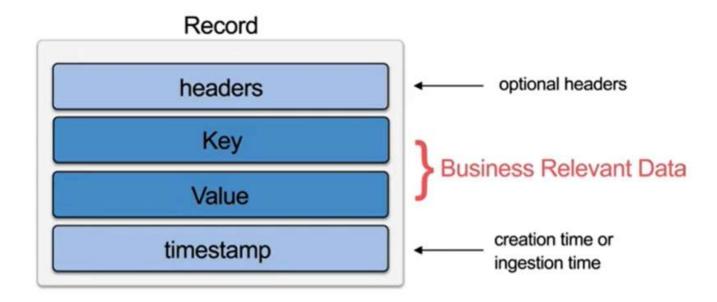
Log Structured Data Flow



The Stream



Data Elements

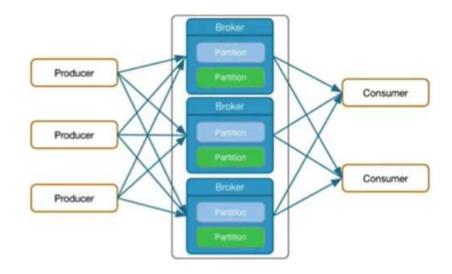


Brokers Manage Partitions

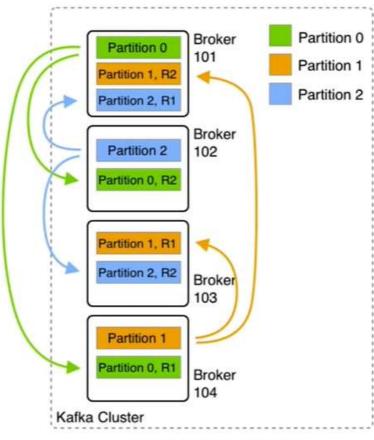
- Messages of Topic spread across Partitions
- Partitions spread across Brokers
- Each Broker handles many Partitions
- Each Partition stored on Broker's disk
- Partition: 1..n log files
- Each message in Log identified by Offset
- Configurable Retention Policy

Broker Basics

- Producer sends Messages to Brokers
- Brokers receive and store Messages
- A Kafka Cluster can have many Brokers
- Each Broker manages multiple Partitions



Broker Replication



Replication (Factor 3)

Producer Basics

- Producers write Data as Messages
- Can be written in any language
 - Native: Java, C/C++, Python, Go, .NET, JMS
 - More Languages by Community
 - REST Proxy for any unsupported Language
- Command Line Producer Tool

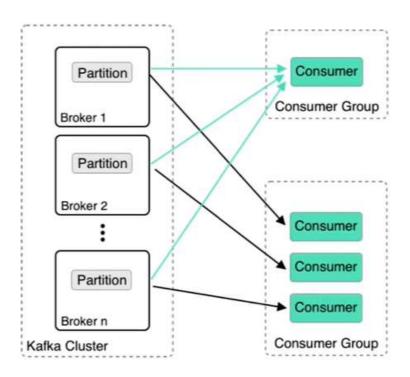
Load Balancing and Semantic Partitioning

- Producers use a Partitioning Strategy to assign each Message to a Partition
- Two Purposes:
 - Load Balancing
 - Semantic Partitioning
- Partitioning Strategy specified by Producer
 - Default Strategy: hash(key) % number_of_partitions
 - No Key → Round-Robin
- Custom Partitioner possible

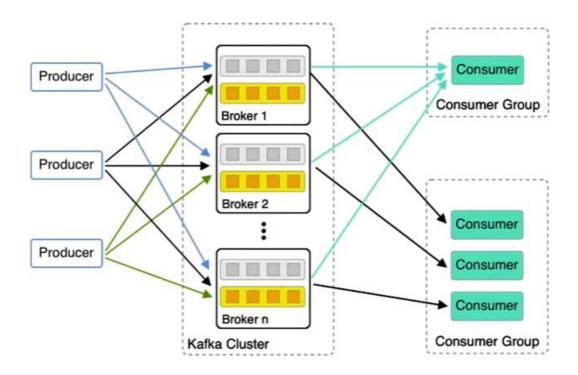
Consumer Basics

- Consumers pull messages from 1...n topics
- New inflowing messages are automatically retrieved
- Consumer offset
 - keeps track of the last message read
 - is stored in special topic
- CLI tools exist to read from cluster

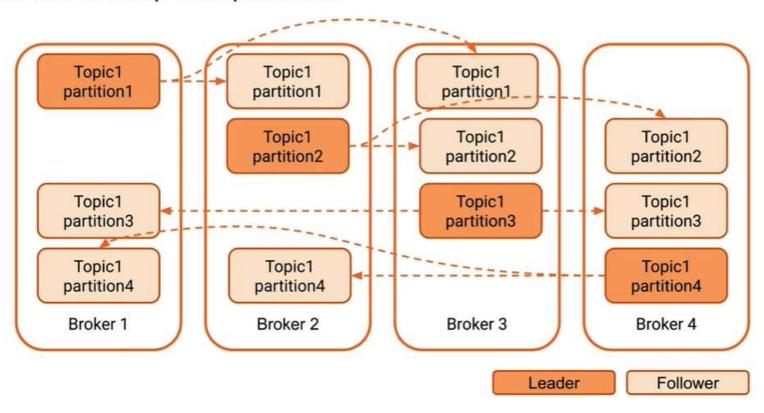
Distributed Consumption



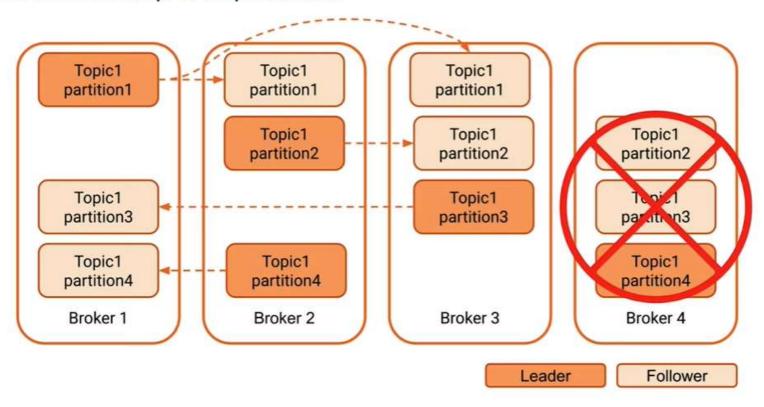
Scalable Data Pipeline



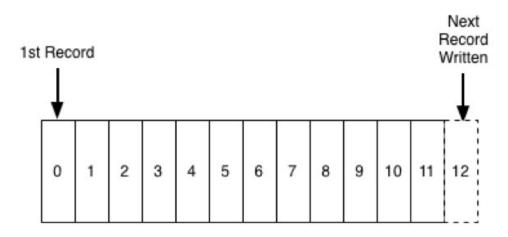
Partition Leadership & Replication



Partition Leadership & Replication



Store data in Kafka?

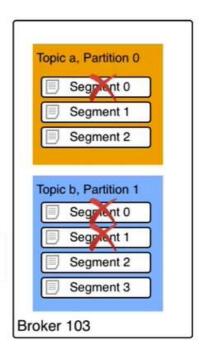


Data Retention Policy

How long do I want or can I store my data?

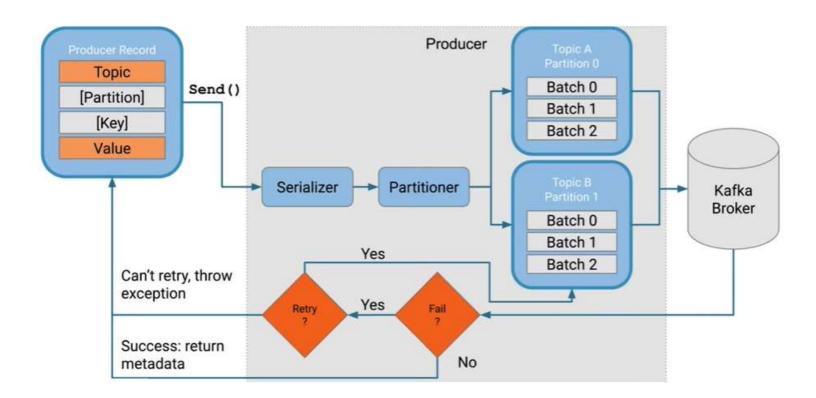
- How long (default: 1 week)
- Set globally or per topic
- Business decision
- Cost factor
- Compliance factor → GDPR

Data purged per segment

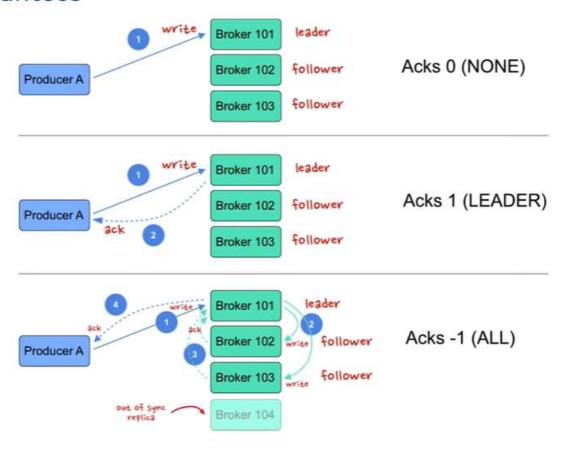


Retention Policy

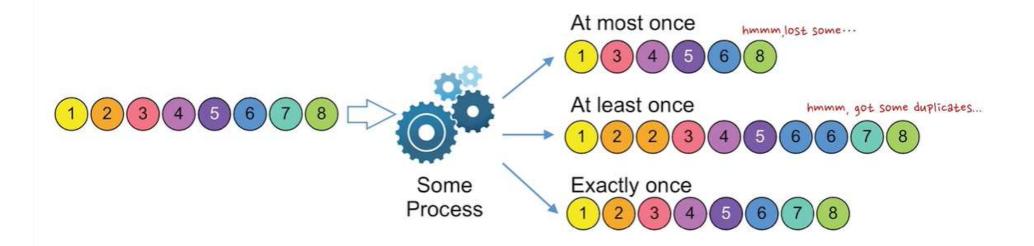
Producer Design



Producer Guarantees



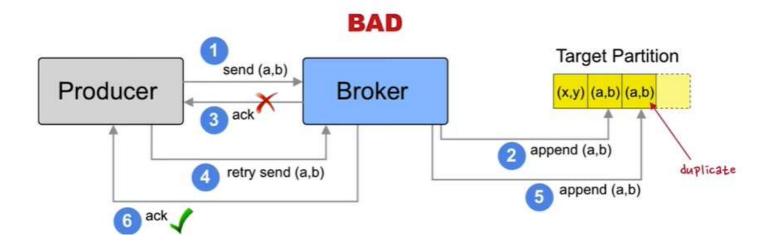
Delivery Guarantees



Idempotent Producers

GOOD





Exactly Once Semantics

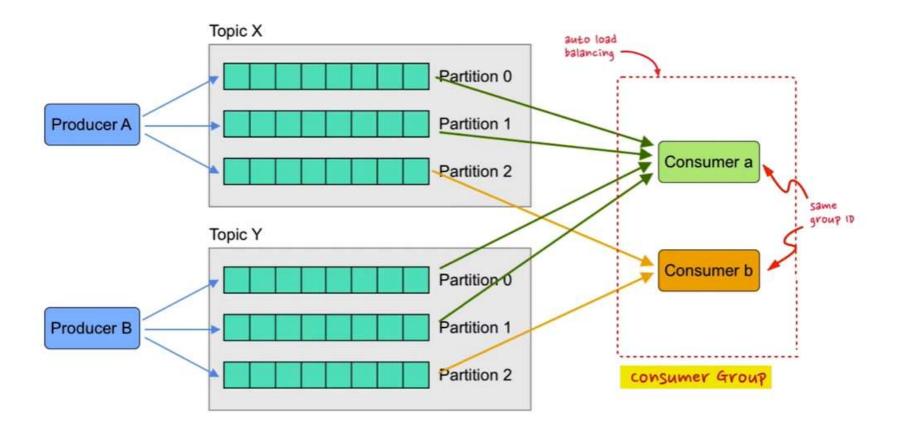
What?

- Strong transactional guarantees for Kafka
- Prevents clients from processing duplicate messages
- Handles failures gracefully

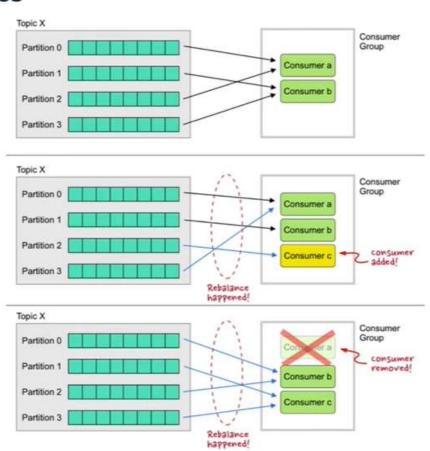
Use Cases

- Tracking ad views
- Processing financial transactions
- Stream processing

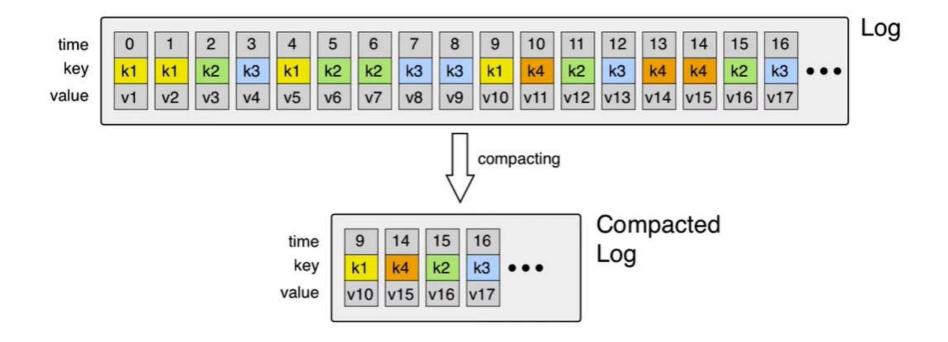
Consumer Groups



Consumer Rebalances

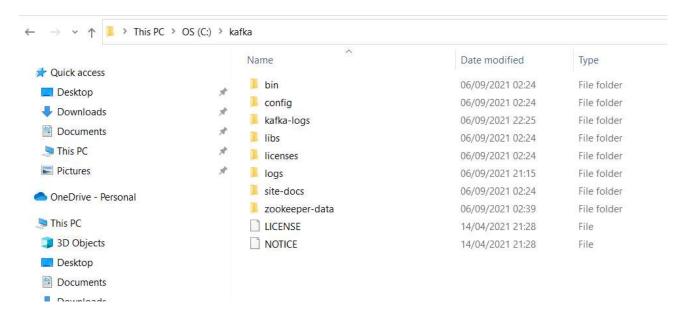


Compacted Topics



Tutorial 1: Kafka installation on Window

- Download Kafka from https://kafka.apache.org/
- Unzip the download file
- Rename the kafka to "kafka" and move it to C:\ drive



Kafka installation on Window

- Open C:\kafka\config\server.properties
- Change the path of log.dir

Kafka installation on Window

- Open C:\kafka\config\zookeeper.properties
- Change the path of dataDir

```
*C:\kafka\config\zookeeper.properties - Notepad++

File Edit Search View Encoding Language Settings Tools Macro Run Plugins Window ?

Server properties 

zookeeper.properties 

the directory where the snapshot is stored.

#dataDir=/tmp/zookeeper

dataDir=C:/kafka/zookeeper-data
```

• By default Apache Kafka will run on port 9092 and Apache Zookeeper will run on port 2181.

Tutorial 2: Run Apache Kafka on Windows

- Start the Kafka cluster
 - Run the following command to start ZooKeeper:

cd C:\kafka\
.\bin\windows\zookeeper-server-start.bat .\config\zookeeper.properties

```
Start Kafka cluster - Start Zookeeper

C:\Users\PC>cd C:\kafka

C:\kafka>.\bin\windows\zookeeper-server-start.bat .\config\zookeeper.properties

[2021-09-06 21:58:44,055] INFO Reading configuration from: .\config\zookeeper.properties (org.apache.zookeeper.server.qu orum.QuorumPeerConfig)

[2021-09-06 21:58:44,064] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
```

Run Apache Kafka on Windows

- Start the Kafka cluster
 - Run the following command to start the Kafka broker:

cd C:\kafka\
.\bin\windows\kafka-server-start.bat .\config\server.properties

```
Start Kafka cluster - Start Kafka broker

C:\Users\PC>cd C:\kafka

C:\kafka>.\bin\windows\kafka-server-start.bat .\config\server.properties

[2021-09-06 22:03:25,045] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.utils.Log4jControllerRegistration$)
```

Run Apache Kafka on Windows

- Produce and consume some messages
 - Run the kafka-topics command to create a Kafka topic named TestTopic

.\bin\windows\kafka-topics.bat --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic TestTopic

Let's create another topic named NewTopic

.\bin\windows\kafka-topics.bat --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic NewTopic

• Let's show list of created topics

.\bin\windows\kafka-topics.bat --list --zookeeper localhost:2181

Run Apache Kafka on Windows

- Produce and consume some messages
 - Run the producer and consumer on separate Command Prompt:

.\bin\windows\kafka-console-producer.bat --broker-list localhost:9092 --topic TestTopic

.\bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic TestTopic --from-beginning

